

Banque Misr internship 2024 Graduation Project Team5

Team members:

- 1- Ahmed Yehia (Scrum Master)
- 2- Marwan Mohammed
- 3- Safia Adel
- 4- Sama Hatem

Part 1: Application Development

We developed a simple Library Management website with (CSS, HTML, and JavaScript) for the frontend and Python Flask for handling the back end. Here is the breakdown of the application features:

- 1- Users for this application are either admins or users.
- 2- The application includes basic routes to get a list of books, retrieve a specific book, add a new book to the library(admin), and borrow and return a book(User).
- 3- Admin Users can add other admins.
- 4- Both admins and users can search for books either with the ISBN or with a keyword that is the title, author, or ISBN.

Here is the login and signup pages:

Library Management System

To exit full screen, press F11

Login

Username

user1

Password

Login

Don't have an account? Sign up

Sign Up

Username

Enter your username

Password

Enter your password

Email

Enter your email

Sign Up

Already have an account? [Login](#)

Here is the admin UI:

Library Management System

Logout

Add Book

Book Title

Author

Select Genre

Publish Year

Add Book

Search by ISBN

Enter ISBN

Search

Search by Keyword

Enter keyword (Title, Author, or ISBN)

Search

Book Inventory

Book Title	Author	ISBN	Genre	Date of Publish	Status	Actions
Sherlock Holmes	Arthur Conan Doyle	1	Mystery	1892	Available	<div>Remove</div>
Death Joke	Agatha Christie	2	Mystery	1930	Available	<div>Remove</div>
Hands on Machine Learning	Aurelien Geron	3	Science	2022	Available	<div>Remove</div>
80 Days around the World	Charles Dickens	4	Fantasy	1873	Available	<div>Remove</div>

Add Book

Search by ISBN

Enter ISBN

Search

Search by Keyword

Enter keyword (Title, Author, or ISBN)

Search

Book Inventory

Book Title	Author	ISBN	Genre	Date of Publish	Status	Actions
Sherlock Holmes	Arthur Conan Doyle	1	Mystery	1892	Available	<div>Remove</div>
Death Joke	Agatha Christie	2	Mystery	1930	Available	<div>Remove</div>
Hands on Machine Learning	Aurelien Geron	3	Science	2022	Available	<div>Remove</div>
80 Days around the World	Charles Dickens	4	Fantasy	1873	Available	<div>Remove</div>
Big Brother	Ahmed Yehia	5	Mystery	12/7/1999	Available	<div>Remove</div>
Big Brother 2	Ahmed Yehia	6	Fantasy	1899	Available	<div>Remove</div>
Book4	Ahmed Yehia	7	Fantasy	1899	Available	<div>Remove</div>


Add new Admin

Username

Password

Add Admin

Here is the User UI:

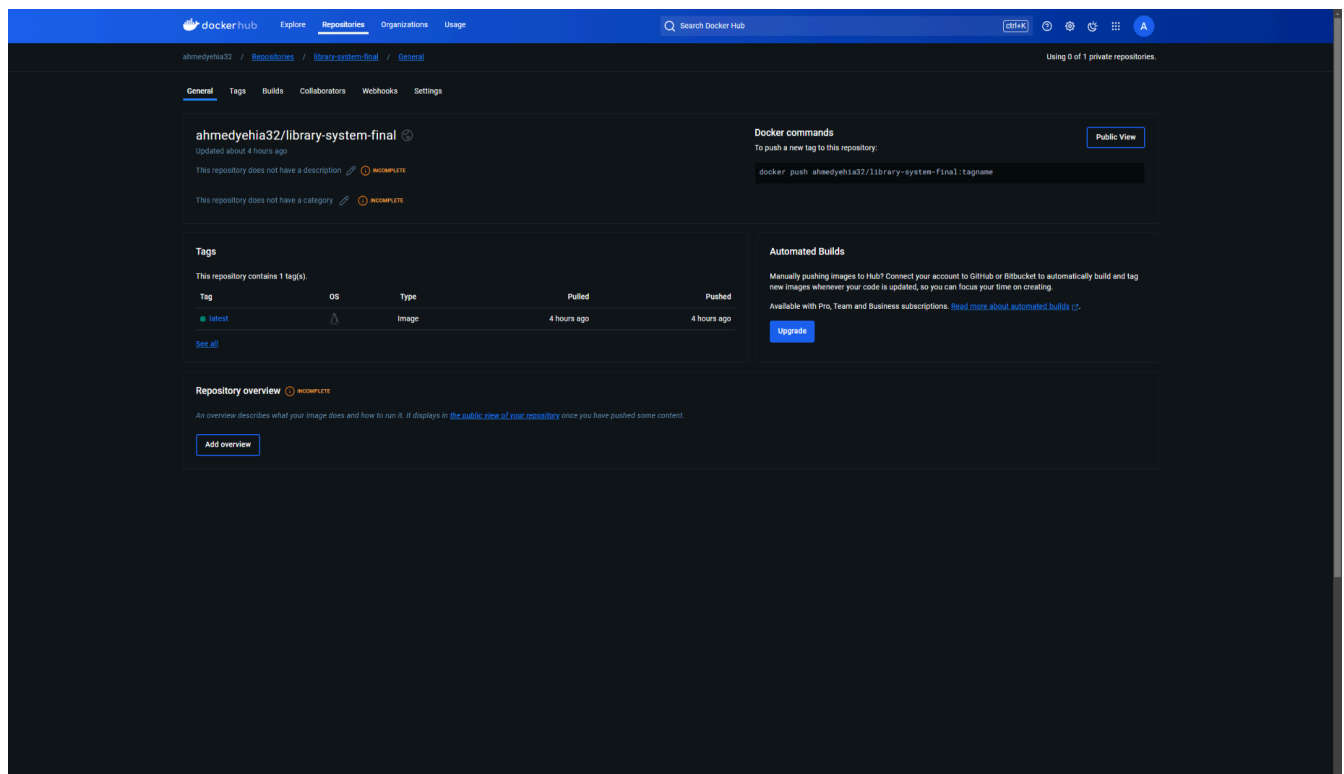


Tools Used:

- 1- Frontend: (JavaScript, HTML, CSS)
- 2- Backend: (Python (Flask))
- 3- Database: our database is Json files.

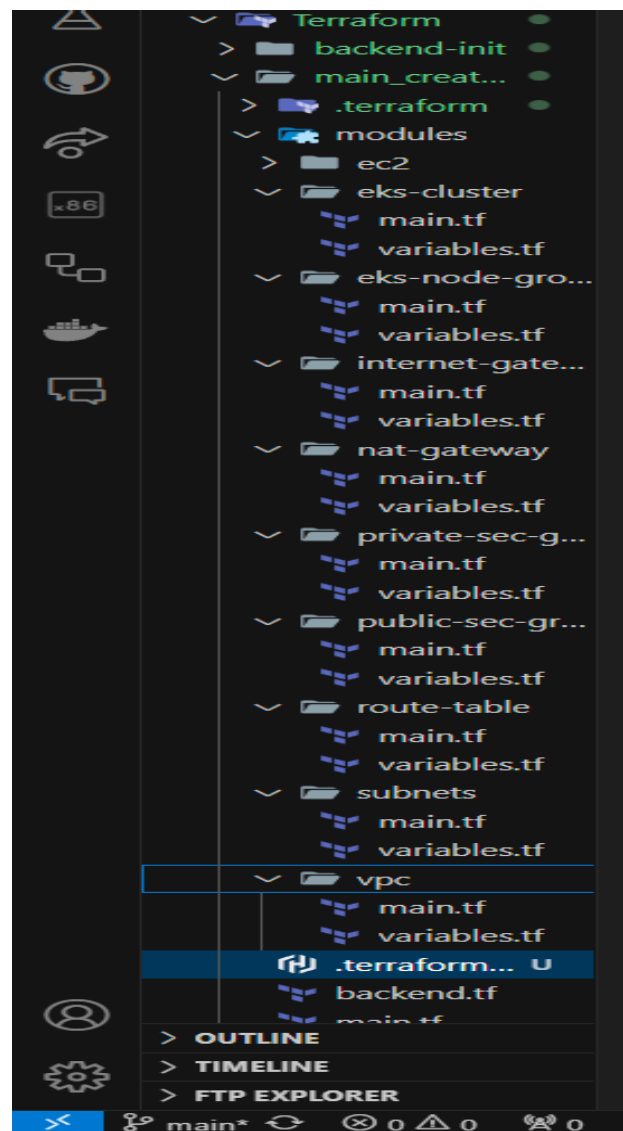
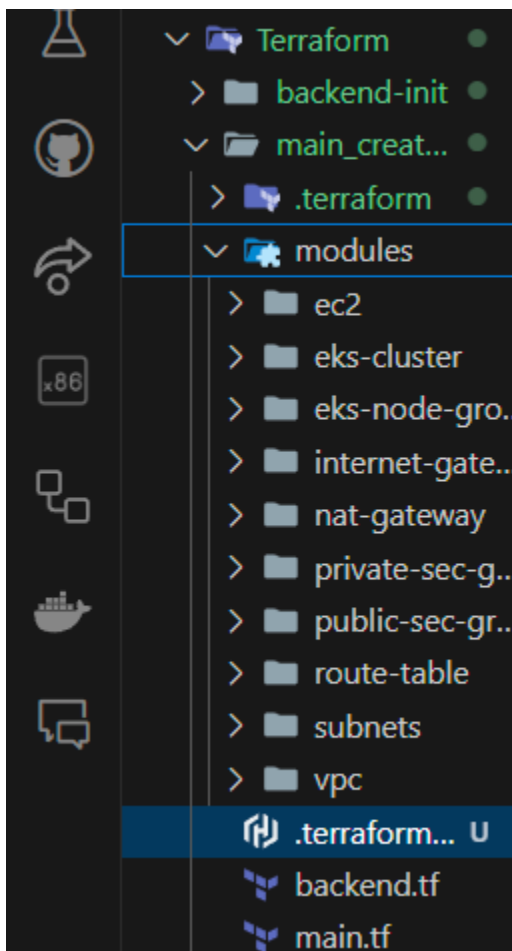
Part 2: Dockerization:

We created a Dockerfile that contains the instructions to build a Docker image. After specifying the base image, setting the working directory, copying the application code, installing dependencies from the (requirements.txt) file, and exposing the necessary port, we containerized our application into a container. After building and testing our Docker Image locally, We pushed the Docker Image to Docker hub.



Part 3: Infrastructure as Code with Terraform

We used Terraform to create our infrastructure on AWS, we had two Terraform modules, one for creating the backend which is the S3 bucket for storing the state file, and the DynamoDB table for preventing multiple changes to the state file at the same time, and the other for creating our main infrastructure(vpc, subnets, Internet gateway, Nat gateway, EKS), our main module contains a module for every AWS resource.



Here is the documentation for each module:

1. VPC Configuration

Module: team5_vpc

Source: ./modules/vpc

Description: Creates a Virtual Private Cloud (VPC) with the specified CIDR block.

Parameters:

cidr_block: 10.0.0.0/16 - The IP address range for the VPC.

name: team5-vpc - Name of the VPC.

2. Subnet Configuration

Modules:

Public Subnet 1

Module: team5_public_subnet1

Source: ./modules/subnets

Parameters:

vpc_id: Refers to the VPC ID created by team5_vpc.

cidr_block: 10.0.1.0/24 - IP address range for the subnet.

availability_zone: eu-central-1a - AWS availability zone.

map_public_ip_on_launch: true - Automatically assign public IP addresses.

name: team5-public-subnet1 - Name of the subnet.

Public Subnet 2

Module: team5_public_subnet2

Source: ./modules/subnets

Parameters:

vpc_id: Refers to the VPC ID created by team5_vpc.

cidr_block: 10.0.2.0/24 - IP address range for the subnet.

availability_zone: eu-central-1b - AWS availability zone.

map_public_ip_on_launch: true - Automatically assign public IP addresses.

name: team5-public-subnet2 - Name of the subnet.

Private Subnet 1

Module: team5_private_subnet1

Source: ./modules/subnets

Parameters:

vpc_id: Refers to the VPC ID created by team5_vpc.

cidr_block: 10.0.3.0/24 - IP address range for the subnet.

availability_zone: eu-central-1a - AWS availability zone.

map_public_ip_on_launch: false - Do not assign public IP addresses.

name: team5-private-subnet1 - Name of the subnet.

Private Subnet 2

Module: team5_private_subnet2

Source: ./modules/subnets

Parameters:

vpc_id: Refers to the VPC ID created by team5_vpc.

cidr_block: 10.0.4.0/24 - IP address range for the subnet.

availability_zone: eu-central-1b - AWS availability zone.

map_public_ip_on_launch: false - Do not assign public IP addresses.

name: team5-private-subnet2 - Name of the subnet.

3. Internet Gateway

Module: team5_internet_gateway

Source: ./modules/internet-gateway

Parameters:

vpc_id: Refers to the VPC ID created by team5_vpc.

name: team5-igw - Name of the Internet Gateway.

4. NAT Gateway

Module: team5_nat_gateway

Source: ./modules/nat-gateway

Parameters:

subnet_id: Refers to the public subnet ID where the NAT Gateway is deployed (team5_public_subnet1).

name: team5-nat-gateway - Name of the NAT Gateway.

5. Route Tables

Modules:

Public Route Table

Module: team5_public_route_table

Source: ./modules/route-table

Parameters:

vpc_id: Refers to the VPC ID created by team5_vpc.

cidr_block: 0.0.0.0/0 - Route all traffic.

gateway_id: Refers to the Internet Gateway ID created by team5_internet_gateway.

name: team5-public-route-table - Name of the route table.

Private Route Table

Module: team5_private_route_table

Source: ./modules/route-table

Parameters:

vpc_id: Refers to the VPC ID created by team5_vpc.

cidr_block: 0.0.0.0/0 - Route all traffic.

nat_gateway_id: Refers to the NAT Gateway ID created by team5_nat_gateway.

name: team5-private-route-table - Name of the route table.

Route Table Associations:

Public Subnet Associations:

team5_public_subnet_a_assoc: Associates team5_public_subnet1 with team5_public_route_table.

team5_public_subnet_b_assoc: Associates team5_public_subnet2 with team5_public_route_table.

Private Subnet Associations:

team5_priv_subnet_a_assoc: Associates team5_private_subnet1 with team5_private_route_table.

team5_priv_subnet_b_assoc: Associates team5_private_subnet2 with team5_private_route_table.

6. IAM Roles for EKS

IAM Role for EKS Cluster

Resource: aws_iam_role.eks_cluster_role

Description: Role for the EKS cluster to assume.

Assume Role Policy:

Allows the EKS service to assume the role.

Policy Attachments:

AmazonEC2FullAccess

AmazonEKSClusterPolicy

AmazonEKSServicePolicy

IAM Role for EKS Node Group

Resource: aws_iam_role.team5_eks_node_role

Description: Role for the EKS node group.

Assume Role Policy:

Allows EC2 instances to assume the role.

Managed Policies:

AmazonEKSWorkerNodePolicy

AmazonEC2ContainerRegistryReadOnly

AmazonEC2ContainerServiceRole

AmazonEKS_CNI_Policy

7. EKS Cluster and Node Group

EKS Cluster

Module: eks_cluster

Source: ./modules/eks-cluster

Parameters:

cluster_name: team5-eks-cluster - Name of the EKS cluster.

cluster_role_arn: ARN of the IAM role for the EKS cluster.

subnet_ids: List of subnet IDs for the cluster.

EKS Node Group

Module: eks_node_group

Source: ./modules/eks-node-group

Parameters:

cluster_name: Name of the EKS cluster (team5-eks-cluster).

node_group_name: team5-node-group - Name of the node group.

node_role_arn: ARN of the IAM role for the node group.

subnet_ids: List of public subnet IDs for the node group.

desired_size: 1 - Desired number of nodes.

max_size: 1 - Maximum number of nodes.

min_size: 1 - Minimum number of nodes.

Part 4: Kubernetes Deployment on EKS:

We deployed our website on Amazon Elastic Kubernetes Service (EKS), after applying our terraform code for creating the infrastructure, we configured our deployment.yaml and service.yaml files for deploying our website on the created EKS cluster.

1- library-management-deployment.yaml file:

```
ent > 📄 library-management-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: library-management-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: library-management
  template:
    metadata:
      labels:
        app: library-management
    spec:
      containers:
        - name: library-management-container
          image: ahmedyehia32/library-system-final
          imagePullPolicy: Always
          ports:
            - containerPort: 5000
            - containerPort: 8000 # Expose the metrics port
```

2- service.yaml file:

```
apiVersion: v1
kind: Service
metadata:
  name: library-management-service
spec:
  selector:
    app: library-management
  ports:
    - name: http
      protocol: TCP
      port: 80
      targetPort: 5000
    - name: metrics
      protocol: TCP
      port: 8000
      targetPort: 8000
  type: LoadBalancer
```

1. Deployment Configuration
(library-management-deployment.yaml)

apiVersion: apps/v1 - The API version used to manage deployments in Kubernetes.

kind: Deployment - Specifies that this YAML defines a Kubernetes Deployment object.

metadata:

name: library-management-deployment - The name assigned to the Deployment object. This name is used to manage and reference the deployment within the cluster.

spec:

replicas: 1 - Defines the number of pod replicas to be created. In this case, one instance of the application is deployed.

selector: This ensures the Deployment targets pods that match specific labels:

matchLabels:

app: library-management - Identifies the pods managed by the deployment based on this label.

template: Defines the specifications for the pods that will be created by the Deployment.

metadata:

labels: Specifies the labels applied to each pod. In this case:
app: library-management

spec: Defines the container configuration for each pod.

containers:

name: library-management-container - The name of the container within the pod.

image: ahmedyehia32/library-system-final - The Docker image to be used. This image contains the library management system application.

imagePullPolicy: Always - Ensures the image is always pulled from the registry, regardless of whether it's cached locally.

Ports:

containerPort: 5000 - The port number inside the container where the application is accessible.

containerPort: 8000:

Purpose: This port is designated for exposing metrics from your application, typically used for Prometheus to scrape.

Usage: If your application has integrated Prometheus metrics (e.g., /metrics endpoint), it would be exposed on this port. Prometheus can then scrape this endpoint to collect metrics data.

2. Service Configuration (service.yaml)

apiVersion: v1 - The version for Kubernetes Service API.

kind: Service - Specifies that this YAML defines a Kubernetes Service object.

Metadata:

name: library-management-service - The name of the service, which will be used to reference this service within the cluster.

Spec:

selector:

app: library-management - This selector matches the labels on the pods created by the deployment, allowing the service to route traffic to them.

Ports:

protocol: TCP - The protocol used by the service to communicate with the pods.

port: 80 - The port that external clients will use to access the service.

targetPort: 5000 - The port on the container where the application is running (as defined in the deployment).

type: LoadBalancer - Specifies that this service is exposed to the internet via a load balancer, making it accessible outside the Kubernetes cluster.

Note: **Port 8000**: The port exposed for metrics collection. It routes traffic directly to port 8000 on the backend pods. This port is used by Prometheus to scrape metrics from the application.

Part 5: CI/CD Pipeline Setup:

Our (CI/CD) pipeline involves automating the process of integrating code changes and deploying our website. We implemented our (CI/CD) pipeline using Jenkins. We configured the pipeline to automatically build, and deploy code whenever the pipeline is executed.

Pipeline Stages:

1. Checkout Code

```
stages {
    stage('Checkout Code') {
        steps {
            git url: "${GITHUB_REPO}", branch: 'main', credentialsId: 'github-token-i
        }
    }
}
```

- **Description:** This stage checks out the code from the specified GitHub repository (`GITHUB_REPO`) using the provided credentials (`github-token-id`). The pipeline will pull the `main` branch.

2- Build Docker Image:

```
stage('Build Docker Image') {  
    steps {  
        script {  
            // Ensure Dockerfile is in the correct directory  
            sh 'docker build -t ${DOCKER_IMAGE_NAME}:${DOCKER_IMAGE_TAG} .'  
        }  
    }  
}
```

- **Description:** In this stage, Jenkins builds a Docker image using the `Dockerfile` located at the root of the project. The image is tagged as `${DOCKER_IMAGE_NAME}:${DOCKER_IMAGE_TAG}`, which is `library-system-final:latest` by default.

3- Push Docker Image:

```

stage('Push Docker Image') {
  steps {
    withCredentials([usernamePassword(credentialsId: 'dockerhub-credentials-id', usernameVariable: 'DOCKER_USERNAME', passwordVariable: 'DOCKER_PASSWORD')]) {
      script {
        // Login to Docker Hub
        sh 'echo $DOCKER_PASSWORD | docker login -u $DOCKER_USERNAME --password-stdin'

        // Tag the Docker image
        sh 'docker tag ${DOCKER_IMAGE_NAME}:${DOCKER_IMAGE_TAG} ${DOCKER_USERNAME}/${DOCKER_IMAGE_NAME}:${DOCKER_IMAGE_TAG}'

        // Push the Docker image and capture output
        sh '''
        set -x
        docker push ${DOCKER_USERNAME}/${DOCKER_IMAGE_NAME}:${DOCKER_IMAGE_TAG} 2>&1 | tee docker_push.log
        '''
      }
    }
  }
}

```

- **Description:**

- Logs in to Docker Hub using credentials (**dockerhub-credentials-id**).
- Tags the built Docker image with the Docker Hub repository name.
- Pushes the Docker image to Docker Hub and logs the output.

4- Deploy to EKS:

```

stage('Deploy to EKS') {
  steps {
    withCredentials([usernamePassword(credentialsId: 'aws-credentials-id', usernameVariable: 'AWS_ACCESS_KEY_ID', passwordVariable: 'AWS_SECRET_ACCESS_KEY')]) {
      script {
        sh 'aws configure set aws_access_key_id $AWS_ACCESS_KEY_ID'
        sh 'aws configure set aws_secret_access_key $AWS_SECRET_ACCESS_KEY'
        sh 'aws configure set region ${AWS_REGION}'
        sh "aws eks --region ${AWS_REGION} update-kubeconfig --name ${EKS_CLUSTER_NAME}"
        sh 'kubectl apply -f Deployment/library-management-deployment.yaml'
        sh 'kubectl apply -f Deployment/service.yaml'
      }
    }
  }
}

```

- **Description:**

- Configures AWS CLI with credentials (`aws-credentials-id`), sets the region, and updates the kubeconfig to connect to the EKS cluster.
- Uses `kubectl apply` to deploy the application in Kubernetes by applying the `library-management-deployment.yaml` and `service.yaml` files.

5- Get the Load Balancer IP:

```
stage('Get Load Balancer IP') {
  steps {
    script {
      sleep(time: 60, unit: 'SECONDS')
      def loadBalancerIP = sh(script: 'kubectl get svc library-management-service -o jsonpath="{.status.loadBalancer.ingress[0].hostname}"', returnStdout: true).trim()
      echo "Load Balancer IP: ${loadBalancerIP}"
    }
  }
}
```

Description:

- Waits for 60 seconds to ensure the Load Balancer is fully initialized.
- Retrieves the Load Balancer IP address or hostname from the `library-management-service` using `kubectl`.
- Displays the Load Balancer IP/hostname in the console output.

Bonus Task: Set Up Monitoring and Logging:

Overview

Prometheus and Grafana are popular tools for monitoring and visualizing metrics. Prometheus is used to collect and store metrics data, while Grafana is used to visualize this data through dashboards.

Components

- 1- Prometheus: A monitoring system and time-series database.
- 2- Grafana: An open-source platform for monitoring and observability, which supports various data sources including Prometheus.

1. Setup Prometheus

1.1 Prometheus ConfigMap

(prometheus-configmap.yaml)

Create a ConfigMap to provide the Prometheus

```
ing > prometheus-configmap.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: prometheus-config
  namespace: monitoring
data:
  prometheus.yml: |
    global:
      scrape_interval: 15s

    scrape_configs:
      - job_name: 'library-management'
        metrics_path: '/metrics'
        static_configs:
          - targets: ['library-management-service.default.svc.cluster.local:8000'] # Fully qualified domain name
```

configuration to the Prometheus deployment.

1.2

Prometheus Deployment

(prometheus-deployment.yaml)

Defines the deployment of Prometheus.

```
ing > prometheus-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: prometheus
  namespace: monitoring
spec:
  replicas: 1
  selector:
    matchLabels:
      app: prometheus
  template:
    metadata:
      labels:
        app: prometheus
    spec:
      containers:
        - name: prometheus
          image: prom/prometheus:latest
          args:
            - "--config.file=/etc/prometheus/prometheus.yml"
          volumeMounts:
            - name: prometheus-config
              mountPath: /etc/prometheus
      volumes:
        - name: prometheus-config
          configMap:
            name: prometheus-config
```

1.3 Prometheus Service (prometheus-service.yaml)

Defines the service for Prometheus.

```
apiVersion: v1
kind: Service
metadata:
  name: prometheus-service
  namespace: monitoring
spec:
  selector:
    app: prometheus
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9090
  type: LoadBalancer
```

1.4 Deployment Commands

Apply the configurations:

```
kubectl apply -f prometheus-configmap.yaml
```

```
kubectl apply -f prometheus-deployment.yaml
```

```
kubectl apply -f prometheus-service.yaml
```

2. Setup Grafana

2.1 Grafana Deployment (grafana-deployment.yaml)

Defines the deployment of Grafana.

```
g > grafana-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: grafana
  namespace: monitoring
spec:
  replicas: 1
  selector:
    matchLabels:
      app: grafana
  template:
    metadata:
      labels:
        app: grafana
    spec:
      containers:
        - name: grafana
          image: grafana/grafana:11.2.0
          ports:
            - containerPort: 3000
          volumeMounts:
            - name: grafana-data
              mountPath: /var/lib/grafana
      volumes:
        - name: grafana-data
          emptyDir: {}
```

2.2 2.2 Grafana Service with LoadBalancer
(grafana-service-loadbalancer.yaml)
Defines the service for Grafana with a LoadBalancer to expose it.

```
apiVersion: v1
kind: Service
metadata:
  name: grafana
  namespace: monitoring
spec:
  selector:
    app.kubernetes.io/name: grafana
    app.kubernetes.io/instance: grafana
  ports:
    - protocol: TCP
      port: 80
      targetPort: 3000
  type: LoadBalancer
```

2.3 Deployment Commands

Apply the configurations:

```
kubectl apply -f grafana-deployment.yaml
```

```
kubectl apply -f grafana-service-loadbalancer.yaml
```

3. Integrating Prometheus with Grafana

3.1 Access Grafana

Get the external IP address of the Grafana service:

```
kubectl get svc grafana -n monitoring
```

Navigate to the Grafana URL provided by the LoadBalancer.

3.2 Add Prometheus as a Data Source

Log in to Grafana (default credentials: admin/admin).

Go to Configuration (Gear icon) > Data Sources.

Add Data Source and select Prometheus.

Configure Prometheus URL to `http://prometheus:9090`

(assuming Prometheus is within the same namespace or adjust accordingly if using an external IP).

3.3 Create Dashboards


Go to Dashboards and click New Dashboard.

Add Panels and configure queries to visualize metrics from Prometheus.

Example Query for HTTP Requests:

```
http_requests_total
```

Some Screenshots from Prometheus:



The screenshot shows the Prometheus Targets page for the `library-management` scrape pool. The table lists the endpoint, state, labels, last scrape time, and scrape duration. The state is `UP`, and the last scrape was 1.782s ago.

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://library-management-service.default.svc.cluster.local:8000/metrics	UP	<code>instance="library-management-service.default.svc.cluster.local:8000"</code> <code>job="library-management"</code>	-1.782s ago	11.40ms	

Some Screenshots from Grafana:

